

USER GUIDE · FIRST-TIME READER EDITION

Macro Narrative Engine

*How to read the market environment, narratives, catalysts,
evidence, and positioning shown on the MNE dashboard.*

Generated **July 16, 2026**

Reflects the dashboard implementation as of this date

Source of truth: current application code and a live persisted run

Evidence-driven

Not a signal generator

Not financial advice

CONTENTS


Table of Contents

01	What MNE Does	The evidence-to-intelligence pipeline
02	The Dashboard in One Page	A labeled reading order
03	How to Use MNE in 60 Seconds	A practical first pass
04	Run Context & Data Quality	What the Trust Summary is telling you
05	Market Environment	What kind of backdrop is present
06	Narrative Leadership	Theme, group, dominance, and rank
07	Pulse	How strong is a narrative's current pull
08	Regime	Is price action confirming the narrative
09	Operating Mode	Macro vs. Nasdaq interpretation
10	Narrative Dynamics	Persistence, acceleration, memory
11	Narrative Brief	The plain-language explanation panel
12	Catalyst Environment	Scheduled events and calendar coverage
13	Market Confirmation & Divergence	Narrative vs. price agreement
14	Breadth	Broad or narrow participation
15	Positioning	Calendar-driven positioning context
16	Market Expression Map	Narrative-to-asset context mapping
17	Trends	Daily leadership history
18	Research Workspace	Investigating a single narrative
19	Historical Replay	Reconstructing a past date
20	What MNE Can & Cannot Tell You	Scope and limits, stated plainly
21	Suggested Daily Workflow	Before, during, and after the session
22	Worked Example	A complete walkthrough of one real run
23	Complete Glossary	Every term, alphabetically

What MNE Does

Macro Narrative Engine (MNE) is an evidence-driven market intelligence system. It helps you understand the current market environment, the dominant narratives driving attention, the catalysts ahead, and — just as importantly — how much you should trust the evidence behind all of it.

Every dashboard run follows the same pipeline. Headlines and data are collected, sorted into narrative themes, measured for strength and change over time, and then layered against real market and calendar context. The result is a single page that tells you what the market conversation currently looks like, and how solid the ground under that read actually is.



MNE IS

An evidence-driven market intelligence system that helps you understand the current market environment, dominant narratives, catalysts, positioning, and the quality of the evidence behind the read.

MNE IS NOT

A stock-picking service, a prediction engine, a trade-signal generator, a guaranteed forecasting tool, or financial advice. MNE helps you understand the environment you are making decisions in — it does not make the decision for you.

The Dashboard in One Page

The live dashboard (the **Overview** page) is built as one long, top-to-bottom scroll. Everything above the fold is primary; everything below is supporting detail you reach for when you want to go deeper. Read it in this order the first few times you use it.

1 Trust Summary & run context
Data Quality banner, selected run, last updated time

START HERE

2 Market environment (hero)
The one-line read of the current backdrop and confidence

PRIMARY

3	Narrative Brief Plain-language explanation of the dominant narrative	PRIMARY
4	What Changed Since Last Run A run-over-run delta of the major moves	SECONDARY
5	Metric grid & Regime Alignment History Regime, Dominant Narrative, Market Environment, Catalysts, Positioning at a glance	PRIMARY
6	Narrative Leadership Ranked cards for the top narrative groups, with Pulse and rotation	PRIMARY
7	Daily Leadership history, Narrative & Group Scores Supporting detail behind the leadership cards	SUPPORTING
8	Market Snapshot QQQ, NVDA, VIX, DXY — the raw price context	SUPPORTING
9	Market, Catalysts, Positioning cards The three environment classifiers, each with its own reasoning	DIAGNOSTIC
10	Top Theme Evidence Example headlines behind the leading themes	SUPPORTING

Three page-level jump links (**Narratives**, **Market**, **Catalysts**, **Trends**) sit in the header for fast navigation once you know the layout. A left-hand rail switches between **Overview** (live dashboard), **Research** (narrative investigation workspace), and **Admin** (operational diagnostics).




Figure 2.1 — The header and Data Quality banner, captured from a live run. This is always the first thing to read.

03 · ORIENTATION

How to Use MNE in 60 Seconds

A practical first pass through the dashboard, in the order you should actually do it.

1. **Confirm the Data Quality banner.** Is the evidence base Strong, Usable, Limited, or Thin? This tells you how much weight to put on everything below it.
2. **Read the Market Environment.** One label (e.g. "Energy Shock," "Clean Risk-On") tells you the backdrop MNE is reading right now.
3. **Identify the dominant narrative.** Check the Narrative Leadership section for the #1 card and its share of total narrative attention.
4. **Check whether it's strengthening or fading.** Look at Pulse, Acceleration, and the rotation badge (Stable, Losing, Challenging, and so on) on that same card.

- 5. **Review upcoming catalysts.** The Catalyst Environment card tells you how dense the next two weeks of scheduled events are.
- 6. **Compare the narrative with market behavior.** Regime Alignment and the narrative/market relationship label tell you whether price action agrees with the story.
- 7. **Open Research Workspace when you need more.** Click "Investigate narrative" on any leadership card to see the underlying evidence.

HOW TO READ THIS GUIDE

Sections 04–17 explain each dashboard element in the order you encounter it on the page. Section 18 covers the Research Workspace, Section 19 covers Historical Replay. If you only read one more section right now, read Section 20 — what MNE can and cannot tell you.

04 · RUN CONTEXT

Run Context & Data Quality

Before you read anything else, the dashboard tells you how much evidence it had to work with, and how current that evidence is.

Which run am I looking at?

MNE writes a new result file every time it runs. The dashboard defaults to the **latest meaningful run** — the newest run that actually produced narrative signal — rather than simply the newest file on disk.

SELECTED RUN	You explicitly chose a run from the Run dropdown. MNE shows exactly that file, with no substitution.
LATEST MEANINGFUL RUN	The default. MNE walks backward from the newest file until it finds one with real theme or group scores, skipping any run that produced no narrative signal. If it had to skip forward, a banner tells you so: <i>"Latest run had no narrative signal; showing latest meaningful run."</i>
LATEST RUN (ADMIN ONLY)	The Admin view shows the literal newest file regardless of content, so an operator can see empty or failed runs the public dashboard hides.

Reading the Trust Summary

The banner at the top states one of five evidence-base labels, a confidence level, a plain-language reason, and a row of supporting facts.

STATE	WHAT IT MEANS
Strong evidence base	Ample accepted evidence, healthy provider network, good coverage.
Usable evidence base	Enough evidence to read with normal confidence, without being exceptional.
Limited evidence base	Meaningful constraints exist in the evidence network — read with more caution.
Thin evidence base	Very little accepted evidence, or the dashboard had to fall back to an older meaningful run.
Data unavailable	No run data or diagnostics could be loaded at all.

The fact row spells out the inputs behind that label: **Accepted evidence** (how many evidence items passed acceptance checks this run), **Providers** (distinct upstream sources), **Evidence network** (a composite network-health read), **Source confidence** (how well sources fetched and contributed this run), and **Coverage** (how broadly that evidence was spread across narratives).

WHAT THIS DOES NOT MEAN

Data confidence describes the evidence base — how much was collected, how fresh it is, how many independent sources agree it exists — not confidence that the market will move in a particular direction. A "Strong evidence base" run can still describe a highly uncertain market; a "Thin evidence base" run can still be directionally correct. They are answering different questions.

05 · PRIMARY READ

Market Environment

Market Environment is the hero headline at the top of the dashboard. It answers one question:

"What kind of market backdrop is currently present?"

MNE checks a fixed, ordered set of conditions — narrative signals (macro stress, AI strength, energy dominance), how concentrated attention is, and two market moves (QQQ and VIX) — and returns the first label that matches:

STATE	ROUGHLY MEANS
Energy Shock	Energy and inflation scores are both elevated alongside high or rising macro stress.
Risk-Off	Volatility sharply higher, Nasdaq down, and Macro Pressure is the dominant narrative.
Fragile Risk-On	Stocks up, but volatility is also sharply higher and macro stress is high — a shaky rally.
AI Momentum	AI / Tech Growth dominates, AI narrative score is strong, and Nasdaq is strongly higher.
Clean Risk-On	Nasdaq strongly higher, volatility flat or down, AI leading, macro stress low.
Macro Stress	Macro stress is high alongside Macro Pressure leadership or rising volatility.
Narrative Rotation	No narrative group clearly dominates; the top two groups are close in score.
Low-Conviction Chop	Nasdaq and volatility both flat, and attention isn't concentrated. Also the fallback state.

Each state also carries a **confidence** (High / Moderate / Low, based on how many corroborating conditions fired) and a **reason** — a fixed, human-written sentence explaining the match.




Figure 5.1 — The hero read from a live run: state "*Broad Macro Read*" is the Operating Mode's plain-language framing (see Section 09); the underlying Market Environment label for this run was *Energy Shock*, visible in the metric grid.

Worked example

On the run captured for this guide, the Market Environment card read **Energy Shock**, with the reason: "*Energy and inflation scores are elevated while macro stress remains high.*" That single label told the user, before reading anything else, that the backdrop was commodity- and inflation-driven rather than a clean risk-on or risk-off tape.

DO NOT INTERPRET IT AS


A forecast of what markets will do next, or an instruction to buy or sell energy-sensitive assets. It is a classification of present conditions, not a projection of future ones.

06 · PRIMARY READ

Narrative Leadership

This section answers: **which story is currently dominating market attention, and by how much?**

Theme, group, and dominance — how they relate





THEME SCORE	A keyword-weighted count of matching headlines for one theme. Stronger keyword matches ("strong" = 3, "medium" = 2, "weak" = 1 point) are summed across non-overlapping matches.
GROUP SCORE	The sum of its constituent themes' scores. Energy / Commodities, for example, sums the <code>energy</code> , <code>oil</code> , and <code>commodities</code> theme scores.
DOMINANT NARRATIVE SHARE	The leading theme's count divided by total mentions across all themes — how much of the total conversation the top story occupies.
CONCENTRATION GAP	The point gap between the top theme and the runner-up. A large gap means a clear leader; a small gap means a close contest.
NARRATIVE CONCENTRATION	The general idea that Dominant Share and Concentration Gap together describe — is attention spread across many stories, or piled onto one?
LEADERSHIP LABEL	Each ranked card also carries a plain-language tag: Dominant Narrative (#1), Challenging Leadership (a close #2), or Secondary Narrative (further back).




Figure 6.1 — Three ranked Narrative Leadership cards from a live run.

- 1 **Rank & leadership label** — #1 "Dominant Narrative," #2 "Challenging Leadership," #3 "Secondary Narrative."
- 2 **Rotation badge** — how this group's rank has moved recently (Stable, Losing, Challenging, Dominant, Holding, Challenged, Emerging). This is a separate signal from Pulse — see Section 17.
- 3 **Pulse, Acceleration, Crowding** — see Section 07 for Pulse; Acceleration and Crowding are covered in Section 10.
- 4 **Mentions and the meter bar** — the group's raw score, shown relative to the #1 group.
- 5 **"Investigate narrative"** — opens this narrative in the Research Workspace (Section 18).

DO NOT INTERPRET IT AS


A prediction that the dominant narrative will keep leading. It is a same-run ranking of existing narrative-group scores, nothing more.

07 · PRIMARY READ

Pulse

WHAT IT IS	A per-narrative-group score (the "strength index") that measures how strongly a narrative is currently pulling attention, weighted mostly by its share of total narrative-group score and how many consecutive runs it has held the lead.
USE IT TO ANSWER	"How much force is behind this narrative right now — is it a passing mention or the center of gravity?"
INPUTS	Group's share of total narrative-group score (weighted most heavily), leadership persistence in consecutive runs, the group's raw current score, and its acceleration state (Accelerating / Stable / Cooling).
WHAT IT DOES NOT MEAN	Pulse is deliberately kept separate from acceleration, crowding, and market direction. A high Pulse does not mean the narrative is accelerating, and it says nothing about whether price is confirming it — check Acceleration (Section 10) and Regime (Section 08) separately.
COMPARE WITH	Narrative Leadership rank (Section 06), Acceleration (Section 10), and Regime Alignment (Section 08) to see whether narrative strength, trend, and market confirmation all agree.

The five Pulse states



Each Pulse reading also carries its own **confidence** (High / Medium / Low), based on how much share and persistence back up the state — a "Strong" Pulse backed by 30%+ share and a multi-run streak is High confidence; a "Strong" reading with thin support is Low confidence.

HOW IT DIFFERS FROM A MARKET SIGNAL

Pulse measures narrative attention and leadership persistence — not price, not volume, not options flow. A Dominant Pulse describes how much the story is dominating coverage, not how the related assets are trading. For that, look at Regime Alignment and the narrative/market relationship state (Sections 08 and 13).

08 · PRIMARY READ

Regime

On the dashboard, "Regime" refers to **Regime Alignment** — a single 0–100 score measuring how well price and market action are *confirming* the current dominant narrative.

MARKET ENVIRONMENT

Classifies the current backdrop (Energy Shock, Risk-Off, Clean Risk-On...) from price moves and narrative signals.

REGIME (REGIME ALIGNMENT)

Scores, 0–100, how well narrative leadership is being confirmed by price, breadth, positioning, and catalysts.

OPERATING MODE

A display-only interpretation lens (Macro or Nasdaq) that turns the same underlying data into a plain-language "read" sentence. It does not change any score.

How the score moves

Regime Alignment starts at a neutral base of 50 and is adjusted by five inputs: the narrative/market relationship state (confirmation adds, divergence or exhaustion subtracts), breadth confirmation, market environment, positioning, and Nasdaq-specific price/volatility context. The result is clamped to 0–100.

SCORE	STATE
80–100	High Alignment

60–79	Moderate Alignment
40–59	Mixed / Low Alignment
20–39	Conflicted Environment
0–19	Highly Conflicted Environment




Figure 8.1 — Regime Alignment History from a live run — the score's day-over-day path, ending at 60.0 (Moderate Alignment).

DO NOT INTERPRET IT AS

A trend-strength or volatility-regime indicator in the traditional technical-analysis sense. It specifically measures narrative-to-market agreement, nothing else.

09 · PRIMARY READ

Operating Mode

Operating Mode is an interpretation lens, not a scoring input. Two modes are currently implemented:

Macro

The default mode. Produces a fixed-format "Broad Macro Read" state and a Moderate-confidence sentence naming the dominant group, market state, and breadth state.

Nasdaq

A richer, Nasdaq-specific read built from QQQ, NVDA, VIX, and DXY moves, macro stress, AI dominance, catalyst positioning, breadth, and Regime Alignment. States include Nasdaq Macro Pressure, Nasdaq Pre-Catalyst Chop, Nasdaq Risk-On Confirmation, Nasdaq Divergence, and Nasdaq Mixed Conditions.

WHAT CHANGING THE MODE AFFECTS

Mode changes only the generated headline sentence and its confidence label — it does not change Pulse, Regime Alignment, group scores, theme scores, or any other underlying calculation. Both modes read the same evidence; they just narrate it differently. An unrecognized mode value silently falls back to Macro.

10 · PRIMARY READ

Narrative Dynamics

A narrative's score tells you where it stands right now. Dynamics tell you where it's headed.

PERSISTENCE

How many consecutive runs a theme or group has held the dominant position. Shown as a run-streak count feeding both Pulse and the Recent Memory panel in Research Workspace.

ACCELERATION

Compares the most recent score change to the prior one. States: **Accelerating** (the move is speeding up), **Cooling** (slowing down), **Stable**, or **Insufficient history** with fewer than three runs on record.

MOMENTUM	A related but separate trend classifier used in reporting and Narrative Memory: Building, Fading, Cooling, Re-accelerating, Flat, or Stable, based on the last two run-over-run deltas.
CROWDING	How concentrated and persistent attention on the dominant narrative has become — High / Moderate / Low , from dominant share, streak length, and acceleration. The engine's own note is explicit: <i>this measures narrative attention crowding, not actual market positioning</i> .

Narrative Memory states

In the Research Workspace, a "Recent Memory" panel classifies a narrative's recent behavior into one label: **Dominant, Persistent, Building, Fading, Recurring, Re-accelerating, Dormant, Emerging**, or **Absent**, alongside plain-language context like appearances, dominance count, streak length, and score/share deltas.

Reading dynamics together — three examples

HIGH SCORE, FADING A group leads on raw score but Acceleration reads Cooling and Momentum reads Fading — the lead is inherited from past runs, not being reinforced today.	LOWER SCORE, ACCELERATING A #3 group with a modest score but Accelerating momentum can be more informative than a flat #1 — it signals where attention is moving next.	DOMINANT BUT LOSING SHARE A group can remain #1 by score while its Dominant Share shrinks run over run — leadership intact, but the margin narrowing.
---	---	--

A NOTE ON "ENTRENCHED" AND "DECELERATING"

These two labels are not currently produced anywhere in MNE's dynamics or memory engine. If you've seen them used informally, treat the closest current equivalents as **Persistent / Dominant** (for "entrenched") and **Cooling** (for "decelerating").

11 · PRIMARY READ

Narrative Brief

The Narrative Brief translates the run's numbers into a short, plain-language explanation — a headline, a summary, and a handful of key points, each traceable back to the exact fields that produced it.




Figure 11.1 — A live Narrative Brief, with its "Data Completeness" badge and key points.

HEADLINE & SUMMARY	Generated from a library of pre-written templates, each gated behind a required-evidence check. A template is only used if every field it needs is actually present in the run.
KEY POINTS	One line per section (Dominant Story, Why It Changed, Market Agreement, Attention, Overall Assessment), each naming the exact modules that supported it.

DATA COMPLETENESS	The label shown on the badge is the brief's internal confidence field (High / Medium / Low), relabeled for the reader. It drops when required inputs for a section are missing, or when more than one generation template conflicted.
WHEN EVIDENCE IS LIMITED	Each section has a fallback template ("Dominant narrative data is unavailable for this run.") that fires automatically rather than fabricating a read. A style guide also blocks hype language — words like "guaranteed," "certain," or "will" are rejected at the template level, before a brief is ever generated.

READ THIS AS

An explanation of the current narrative structure, not as a forecast. Every sentence in the brief is grounded in a specific, logged field from this run — it is describing what already happened in the evidence, not projecting what happens next.

12 · PRIMARY READ

Catalyst Environment

The Catalyst Environment answers: **how much scheduled market-moving news is coming, and how well-verified is that calendar?**

Density and state

MNE counts upcoming **red** (high-impact) and **orange** (medium-impact) events over the next 14 days. Density score = (red events × 2) + orange events.

DENSITY SCORE	STATE
0	Quiet
1	Light
2–3	Moderate
4–5	Elevated
6+	Heavy

What counts as red vs. orange

High-impact (red) events currently include CPI, PPI, the Employment Situation (the monthly jobs report), FOMC decisions and press conferences, GDP, and Personal Income and Outlays (which carries the PCE price index). Medium-impact (orange) events include JOLTS, the Employment Cost Index, Import/Export Price Indexes, FOMC Minutes, International Trade, and the EIA's weekly petroleum and natural-gas reports.

IMPORTANCE IS NOT DIRECTION

A CPI print being flagged "high-impact" says nothing about whether it will be bullish or bearish. The red/orange label is purely about expected attention and volatility, never about direction. Nothing in the catalyst engine assigns bullish or bearish meaning to any event.

Verified quiet vs. calendar unavailable

A density score of 0 alone is not enough for MNE to claim "nothing is coming." It only labels the calendar **Quiet** when the automatic macro-calendar refresh is healthy and current. If that refresh was partial, stale, missing, or unparseable, MNE reports **Calendar Unavailable** or **Calendar Stale** instead of Quiet — and lowers catalyst confidence — so a failed check is never mistaken for a genuinely clear calendar.

VERIFIED QUIET	CALENDAR UNAVAILABLE / STALE
Macro calendar refreshed successfully, and no red/orange events found.	The refresh itself failed, was partial, or every stored event date is in the past — MNE cannot confirm there's nothing coming.




Figure 12.1 — Live Catalysts card showing state *Heavy* with a Calendar Notice — the calendar found events, but the refresh itself was only partial, so MNE flags that upcoming coverage may be incomplete even while reporting a real state.

Company earnings catalysts are sourced automatically for a fixed watchlist of ten large-cap names (NVDA, AAPL, MSFT, GOOGL, META, AMZN, AMD, TSM, AVGO, TSLA); anything outside that list must be added manually to the catalyst calendar.

13 · PRIMARY READ

Market Confirmation & Divergence

This is the narrative/market relationship classifier — it compares what the narrative is saying against what price is actually doing.

STATE	WHAT IT MEANS
Narrative Confirmation	The AI / Tech Growth narrative is stable or accelerating, and QQQ and NVDA are both up while volatility stays controlled — price is validating the story.
Narrative Divergence	The AI / Tech Growth narrative is still strong by score, but QQQ or NVDA is down — the story isn't being confirmed by price.
Narrative Exhaustion	Attention crowding is High, the same group has led for 5+ runs, share is 50%+, but price is now weakening — the narrative may be overextended.
Narrative Rotation	The dominant group changed from the prior run, or the leader is cooling while another group accelerates.
Neutral / Mixed	No clear pattern, or required inputs were unavailable.

A SCOPE NOTE

The Confirmation and Divergence states are currently only defined for the AI / Tech Growth group. Other dominant groups (Macro Pressure, Energy / Commodities, Geopolitical Risk) will show Rotation or Neutral / Mixed instead — they don't yet have their own confirmation logic.

Examples

- AI narrative strengthens, and QQQ and NVDA both confirm with volatility contained → **Narrative Confirmation**.
- Energy headlines accelerate, but energy-sensitive assets don't respond → this shows up as **Rotation** or **Neutral / Mixed**, since the confirmation logic only tracks AI / Tech Growth directly.
- Macro stress rises while defensive positioning firms up → reflected through Regime Alignment's market-environment and positioning inputs, not this classifier directly.

DO NOT INTERPRET IT AS

A trade signal. "Confirmation" means price and narrative currently agree — it is not an instruction to buy, and "Divergence" is not an instruction to sell.

14 · PRIMARY READ

Breadth

Breadth checks whether index-level moves are broadly shared or concentrated in a handful of large names, using SPY, RSP (equal-weight S&P 500), QQQE (equal-weight Nasdaq-100), and IWM (small caps) alongside QQQ, NVDA, and VIX.

BROAD PARTICIPATION

QQQ, SPY, and RSP all rising together, volatility flat or down — strength shared across cap-weighted, equal-weight, and small-cap names.

NARROW PARTICIPATION

QQQ and NVDA up while equal-weight and small-cap proxies lag or fall — the index move is a small number of names doing the work.

STATE	MEANING
Strong Breadth Confirmation	Broad-based strength across cap-weighted, equal-weight, and small-cap participation.
Broad Risk-Off	Broad-based selling across the same participation set, with volatility rising.
Narrow Leadership	Headline strength concentrated in QQQ/NVDA while broader participation lags.
Weak Breadth	Headline index holding up while underlying participation is deteriorating.
Neutral Breadth	No clear broad confirmation or deterioration — the fallback state.

WHAT BREADTH DOES NOT PROVE

Breadth can support or weaken confidence in a narrative read, but it does not by itself prove a move is real or fake — it is one more piece of corroborating context, best read alongside Regime Alignment and the narrative/market relationship state.

15 · PRIMARY READ

Positioning

On the current dashboard, "Positioning Environment" is a **calendar-driven** classifier — it describes how close the market is to a high-impact scheduled event, not trader positioning, options skew, or futures data.

STATE	TRIGGER
Active Catalyst Positioning	A high-impact event is scheduled today.
Heavy Event Positioning	Catalyst density score is 6 or higher.
Elevated Positioning Activity	Catalyst density score is 4 or higher.
Pre-Catalyst Positioning	A high-impact event is scheduled tomorrow.
Light Event Positioning	Some catalyst density, below the Elevated threshold.
Normal Positioning Environment	No nearby high-impact catalysts.
Unknown	Catalyst data wasn't available to classify against.




Figure 15.1 — In this live run, Positioning read *Unknown* — "No catalyst environment data available." — a reminder that this classifier depends entirely on the catalyst calendar being present in the shape it expects.

WHAT THIS DOES NOT MEAN

MNE does not currently compute trader positioning, options-market crowding, one-sided futures positioning, or unwind risk from market data — those concepts are not implemented in the codebase. What's labeled "Positioning" here is contextual and calendar-based: it tells you how close you are to a scheduled catalyst, not how crowded a trade is. Do not use it as a timing signal for entries or exits. The separate "narrative crowding" concept (Section 10) measures attention concentration, not market positioning, either.

16 · CONTEXT MAPPING

Market Expression Map

MNE maintains a static map from each theme to the assets typically associated with it — primary names, secondary names, and offsets (hedges or inversely-related instruments).

NOT CURRENTLY SHOWN ON THE DASHBOARD

This mapping is computed on every run and saved into the run's data, but it is **not rendered anywhere in the current dashboard interface**. It exists in the underlying data and in text/CLI reports, and it can be cited as supporting context inside a Narrative Brief, but there is no dashboard panel for it today. If you see asset tickers referenced in brief text, this map is the likely source.

When it is surfaced (in reports or future dashboard views), each entry is explicitly labeled: *"These assets are narrative-expression proxies and not trade recommendations."*

IF THIS IS EVER ADDED TO THE DASHBOARD

The Market Expression Map identifies assets connected to a narrative. It does not tell the user to buy or sell them.

17 · SUPPORTING DETAIL

Trends

The Narrative Leadership History table tracks which group led each day, its dominant share, and the gap to the runner-up — the run-over-run story behind the leadership cards.

NARRATIVE LEADERSHIP HISTORY				Leadership shifted from Macro Pressure to Energy / Commodities.		
Daily Leadership						
Day	Dominant Group	Theme	Dominant Share	Gap		
07/03	AI / Tech Growth 28 score	ai		58.3%		16
07/06	AI / Tech Growth 25 score	ai		55.6%		11
07/07	Energy / Commodities 2 score	energy		66.7%		1
07/08	Energy / Commodities 19 score	energy		51.3%		11
07/09	Energy / Commodities 15 score	energy		38.5%		5
07/11	Energy / Commodities 15 score	energy		41.7%		4
07/13	AI / Tech Growth 13 score	ai		38.2%		0
07/14	Macro Pressure 42 score	inflation		39.6%		9
07/15	Macro Pressure 26 score	inflation		27.3%		0
07/16	Energy / Commodities 19 score	energy		45.2%		9

Figure 17.1 — Ten trading days of daily leadership history from a live run. Notice the leadership change from AI / Tech Growth to Energy / Commodities around 07/07, and back-and-forth with Macro Pressure through mid-July.

One run per day, not one row per run

MNE can be run more than once on the same day. To keep this history meaningful, the daily trend views (this table, and the CLI trend reports) collapse each calendar day down to its single most recent run before charting it — so re-running MNE five times on a Tuesday does not inflate Tuesday's weight in the trend line to five data points. This day-collapsing logic is separate from the "Recent Memory" panel in Research Workspace, which instead looks at the last N runs directly, regardless of how many fell on the same day.

HOW TO READ THE GAP COLUMN

Gap is the point difference between the day's leader and the runner-up. A widening gap over successive days means the leading narrative is pulling away; a narrowing gap means the field is tightening — worth cross-checking against the rotation badge on the live leadership card.

18 · GOING DEEPER

Research Workspace

When the dashboard's summary isn't enough, click **Investigate narrative** on any leadership card — or open **Research** from the left rail — to enter a focused evidence review for a single narrative.

The screenshot shows the 'Research Workspace' interface. On the left is a sidebar with 'Overview', 'Research' (highlighted), and 'Admin'. The main content area is titled 'RESEARCH WORKSPACE' and 'Choose a narrative to investigate'. Below this is a grid of narrative cards. Each card has a title, a score in a green bubble, and an 'Investigate narrative' button. The cards are:

- Energy / Commodities** (Score: 19) - GROUP NARRATIVE
- Macro Pressure** (Score: 13) - GROUP NARRATIVE
- AI / Tech Growth** (Score: 10) - GROUP NARRATIVE
- Energy** (Score: 19) - THEME SIGNAL
- Ai** (Score: 10) - THEME SIGNAL
- Inflation** (Score: 9) - THEME SIGNAL
- Recession** (Score: 4) - THEME SIGNAL

Each card also includes the text 'Investigation candidate from the selected run.' and 'Investigate narrative'.

Figure 18.1 — The narrative selector: one card per investigable narrative in the selected run, both at the group level and the individual theme level.

The four zones of a narrative investigation

- 1 Snapshot**
"What is happening?" — score, Pulse, Regime Alignment, evidence breadth

2

Explanation

"Why is it happening?" — the narrative-scoped portion of the Narrative Brief

3

Evidence

"What evidence supports it?" — Supporting Evidence, the Evidence Reader, Coverage, and Source Summary

4

Research

"Where should I look next?" — links back to the dashboard, admin diagnostics, and related catalysts




Figure 18.2 — Supporting Evidence list (left) with source, provider, and timestamp per item, and the Evidence Reader panel (right) for a quick, headline-level summary of a selected item.

HEADLINE-LEVEL EVIDENCE, STATED HONESTLY

The Evidence Reader panel is explicit about its own limits. Every item includes the line: *"This interpretation is based on the persisted headline and evidence metadata. Full article text is not available in MNE."* and *"This is not a complete article summary."* MNE tells you directly when its knowledge of a story stops at the headline.

Coverage in this view (Minimal / Limited / Moderate / Broad / Extensive) describes how broad the supporting evidence is for one narrative — how many items, sources, and providers back it — and is explicitly not a judgment of narrative quality. **Source Summary** lists each contributing source with its health and freshness state.

19 · GOING DEEPER (ADMIN)**Historical Replay**

Historical Replay reconstructs what MNE would have shown on a past date, using only evidence that MNE had actually stored at or before that moment. It is an admin-only, read-only feature.

CUTOFF

Evidence published
on or before the cutoff
→ included

Evidence published
after the cutoff
→ excluded

Why replay results may be thin

Two separate limits apply. First, replay can only draw on evidence MNE actually persisted at the time — if live collection wasn't running yet for a given date, there is nothing to replay. Second, an optional **Historical Backfill** can add older evidence from outside sources, but today that backfill covers exactly one source: Federal Reserve / FOMC public communications. It is explicitly labeled partial coverage, not a full market-news reconstruction.

How replay differs from current intelligence

	LIVE DASHBOARD	HISTORICAL REPLAY
Evidence	Current, ongoing collection	Persisted evidence only, up to a chosen cutoff
Live fetching	Yes	Never
Access	Public dashboard	Admin only, read only
Use case	What's happening now	What MNE could have known on a past date

Historical Comparison takes two completed replay artifacts and shows what changed between them — dominant theme/group shifts, score deltas, and evidence-base differences — without generating any new evidence or scores.

READ WITH CAUTION WHEN EVIDENCE IS THIN

Historical Research views carry their own warnings when source coverage is limited, and state directly: *"Historical coverage may be incomplete if MNE had limited stored evidence for this date."* Older replays should be trusted less by default, not more.

20 · SCOPE

What MNE Can & Cannot Tell You

MNE can help with

- Identifying which narrative is dominating market attention right now
- Understanding how much and how reliable the underlying evidence is
- Recognizing when narrative leadership is changing or rotating
- Seeing which high-impact events are coming up
- Comparing a narrative's strength against how price is actually behaving
- Structuring how you prepare for a session, methodically and repeatably

MNE cannot guarantee

- Market direction, on any timeframe
- Exact price targets
- Profitable trades
- The outcome of any scheduled event
- Timing precision for entries or exits
- Complete historical evidence coverage
- Perfect or continuous source availability

THE ONE SENTENCE TO REMEMBER

MNE helps a user understand the environment in which they are making decisions. It does not make the decision for them.

21 · PRACTICE

Suggested Daily Workflow

Before market open

- Check run freshness and the Data Quality banner — confirm you're not looking at a stale or thin run.
- Read the Market Environment and identify the dominant narrative.
- Review upcoming catalysts and note anything red (high-impact) in the next day or two.
- Note the confirmation/divergence state — is price already agreeing with the story, or not yet?

- Ask yourself: what specific evidence, if it appeared today, would invalidate this read?

During the session

- Use MNE as context for interpreting what you're seeing — not as an entry trigger by itself.
- Watch whether price action is confirming or rejecting the environment MNE described this morning.
- Resist forcing a trade just because a narrative is dominant — dominance is a description of attention, not permission to act.

After the session

- Compare what changed against the "What Changed Since Last Run" card.
- Review acceleration and persistence — is the story building, cooling, or unchanged?
- If something surprised you, open Research Workspace and read the actual supporting evidence, not just the score.

22 · COMPLETE WALKTHROUGH

Worked Example

A full read-through of one real, live MNE run, captured for this guide on **2026-07-16 at 11:57**. Every figure below is sourced directly from that run.

DATA QUALITY Limited evidence base High confidence · 91 accepted evidence · 13 providers	MARKET ENVIRONMENT Energy Shock Energy and inflation elevated, macro stress high	DOMINANT NARRATIVE Energy / Commodities 19 mentions · 45.2% dominant share · gap of 9
PULSE Strong High confidence · Accelerating	REGIME ALIGNMENT 60 / 100 Moderate Alignment	NARRATIVE / MARKET RELATIONSHIP Narrative Rotation Leadership changed from Macro Pressure this run
CATALYST ENVIRONMENT Heavy Macro calendar refresh was partial	BREADTH Neutral Breadth No clear broad confirmation either way	POSITIONING Unknown No catalyst environment data available for this classifier

Walking through it

Data quality first. The banner read "Limited evidence base" at High summary confidence — enough evidence to read the run, with a stated network constraint. That set the caution level for everything that followed.

Environment and dominant narrative. Market Environment read Energy Shock. Narrative Leadership confirmed Energy / Commodities as the clear #1, having just taken leadership from Macro Pressure — visible directly in the "What Changed Since Last Run" card ("Energy / Commodities took leadership from Macro Pressure").

Pulse and Regime. Pulse was Strong with High confidence and an Accelerating trend — this wasn't a narrative merely holding on to a prior lead, it was actively strengthening. Regime Alignment sat at 60 (Moderate Alignment) — price and market context were partially, not fully, confirming the story.

Persistence and catalysts. Research Workspace showed Energy / Commodities dominant in 10 of the last 10 meaningful runs, with a 10-run streak and +18.0 point share delta — a narrative gaining ground, not just persisting. The Catalyst Environment read Heavy, but with a Calendar Notice flagging a partial macro-calendar refresh, so upcoming coverage might understate what's actually scheduled.

Confirmation and positioning. The narrative/market relationship read Narrative Rotation (a same-run leadership change), and Breadth was Neutral — no strong broad confirmation either direction. Positioning read Unknown, because the calendar data it depends on wasn't available in the shape that classifier needed.




Figure 22.1 — The Recent Memory panel for Energy / Commodities from this run.

WHAT THE DASHBOARD SUGGESTS

Energy / Commodities was the clear, strengthening center of market attention, backed by a real and accelerating evidence trend — not a one-run spike.

WHAT REMAINS UNCERTAIN

Regime Alignment was only Moderate, breadth was Neutral, and positioning context was unavailable — market confirmation was partial, not complete, and the catalyst calendar itself was flagged as an incomplete refresh.

WHAT THE USER SHOULD MONITOR NEXT

Whether Regime Alignment moves toward High Alignment as more market confirmation arrives, whether the macro calendar refresh completes and changes the catalyst read, and whether Energy / Commodities' Pulse and share continue to build or begin cooling in the next run.

This walkthrough is a description of one run's data, not a trade recommendation. See Section 20 for what MNE can and cannot tell you.

Complete Glossary

Every term used on the dashboard, defined in one place. Where a state applies specific numeric thresholds, see the relevant numbered section for detail.

Accepted Evidence	An evidence item that passed acceptance checks this run (not a duplicate, not stale, has a parseable timestamp). Counted on the Trust Summary banner.
Acceleration	Whether a narrative's score change is speeding up (Accelerating), slowing (Cooling), or flat (Stable), based on the last three runs. §10
Breadth	Whether index-level strength or weakness is shared broadly across cap-weighted, equal-weight, and small-cap proxies, or concentrated in a few names. §14
Catalyst	A scheduled, market-relevant event — a macro data release, a Fed decision, an earnings report — tagged as high-impact (red) or medium-impact (orange). §12
Catalyst Density	A weighted count of upcoming red and orange catalysts within a 14-day lookahead window. §12
Catalyst Environment	The dashboard card combining catalyst density, state, and calendar-coverage confidence. §12
Concentration Gap	The score gap between the top theme and its closest rival. Larger = clearer leader. §06
Confidence	Appears throughout the dashboard with different scopes each time — evidence-base confidence, catalyst confidence, Pulse confidence, brief confidence. Always read it next to the label it's attached to, not on its own.
Confirmation	"Narrative Confirmation" — price action (currently tracked for AI / Tech Growth) agreeing with a strengthening narrative. §13
Coverage	How broadly a narrative's supporting evidence spans across items, sources, and providers: Minimal, Limited, Moderate, Broad, or Extensive. Not a quality judgment. §18
Crowding	How concentrated and persistent narrative <i>attention</i> has become (High/Moderate/Low) — explicitly not a measure of market positioning. §10, §15
Data Completeness	The user-facing label for a Narrative Brief's internal confidence score (High/Medium/Low). §11
Divergence	"Narrative Divergence" — a strong or accelerating narrative (currently tracked for AI / Tech Growth) whose related assets are falling instead of confirming it. §13
Dominant Narrative	The single theme with the highest score this run. §06
Dominant Narrative Group	The narrative group (of four fixed groups) with the highest summed score this run. §06
Dominant Narrative Share	The dominant theme's mention count divided by total mentions across all themes. §06
Energy Dominance	A narrative-strength signal for the energy theme: HIGH, ELEVATED, or LOW, based on share and concentration. Feeds Market Environment.
Entrenched	Not currently produced by MNE's dynamics or memory engine. The closest current states are Persistent and Dominant. §10
Evidence Funnel	The count of items fetched vs. accepted (fresh) vs. rejected (stale, duplicate, or unknown-timestamp) each run.
Evidence Network	The Trust Summary fact label for the composite Network Confidence / Network Health read. §04
Evidence Reader	The Research Workspace panel that gives a headline-level summary of one evidence item, with explicit "what MNE knows / does not know" sections. §18
Exhaustion	"Narrative Exhaustion" — high, persistent attention crowding on a dominant narrative even as price begins weakening. §13

Feed Health	Per-source fetch outcome this run: Healthy, Partial, Empty, Offline, Redirected, Parse Error, Rate Limited, Blocked, or Unknown.
Freshness	Whether evidence or a source's latest item arrived within its configured recency threshold: Fresh, Stale, Unknown, or (source-level) Quiet.
Group Score	The sum of a narrative group's constituent theme scores. §06
Historical Replay	An admin-only reconstruction of what MNE would have shown on a past date, using only evidence persisted at or before a chosen cutoff. §19
Latest Meaningful Run	The newest result file that actually produced narrative signal — the dashboard's default, skipping empty or failed newer runs. §04
Macro Stress	A composite signal (HIGH/ELEVATED/MODERATE/LOW) built from energy, rates, inflation, and recession theme intensity. Feeds Market Environment and Regime.
Market Environment	The classified current market backdrop (Energy Shock, Risk-Off, Clean Risk-On, etc.). §05
Market Expression	A static per-theme map of associated primary, secondary, and offset tickers. Computed, but not currently shown on the dashboard. §16
Momentum	A trend classifier (Building/Fading/Cooling/Re-accelerating/Flat/Stable) used in reporting and Narrative Memory. §10
Narrative	A recurring market story tracked at either theme level (e.g. "energy") or group level (e.g. "Energy / Commodities"). §06
Narrative Brief	The generated plain-language explanation of the current dominant narrative and its supporting evidence. §11
Narrative Concentration	How much of total narrative attention is piled onto the leading story, vs. spread across several. §06
Narrative Group	One of four fixed roll-ups: AI / Tech Growth, Macro Pressure, Energy / Commodities, Geopolitical Risk. §06
Narrative Leadership	The ranked view of narrative groups by score, with rotation and Pulse context. §06
Narrative Persistence	Consecutive runs a theme or group has held the dominant position. §10
Network Confidence	A composite trust signal (HIGH/MODERATE/LOW/VERY_LOW/UNKNOWN) built from network health, source confidence, source reliability, and coverage together. §04
Network Health	Feed-infrastructure status by provider and category: Healthy, Partial, Degraded, or Critical. §04
Operating Mode	The display lens (Macro or Nasdaq) used to narrate the same underlying data. Does not change scoring. §09
Partial Coverage	Not a single fixed state — used loosely for two distinct things: a partial macro-calendar refresh (§12), and Minimal-level narrative evidence coverage, labeled "Thin" on the Trust Summary (§04).
Positioning	A calendar-proximity classifier — not trader positioning, options data, or unwind risk. §15
Pulse	A per-group strength index (Dormant/Emerging/Building/Strong/Dominant), weighted mostly by share and leadership persistence. §07
Regime	Regime Alignment — a 0–100 score measuring how well price action confirms the dominant narrative. §08
Rotation	Two distinct uses: "Narrative Rotation" (a narrative/market relationship state, §13) and a leadership card's rotation badge (Dominant/Holding/Challenged/Challenging/Emerging/Losing/Stable, describing rank stability, §06).
Source Confidence	A single-run, whole-network read of fetch success, contribution, and freshness ratios: High/Moderate/Low/Unknown. §04
Source Reliability	A per-source, multi-run behavioral track record: Stable/Watch/Repeated Failure/Quarantine Recommended/Unknown. §04

Theme	The most granular narrative unit — e.g. energy, ai, inflation, recession, rates. §06
Theme Score	A keyword-weighted count of matching headlines for one theme. §06
Thin Evidence Base	A Trust Summary state: very little accepted evidence this run, or a forced fallback to an older meaningful run. §04
Trust Summary	The Data Quality banner at the top of the dashboard, stating evidence-base state, confidence, and supporting facts. §04
Verified Quiet Calendar	A "Quiet" catalyst state backed by a successful, current macro-calendar refresh — distinct from a calendar that's simply unavailable. §12